

# Develop Generative AI Applications Get Started

# Artificial Intelligence

## AI models

- Learn from massive existing **data** through **training**



Discriminative AI

Generative AI

## Discriminative AI

Distinguishes between different classes of data

Limitations:

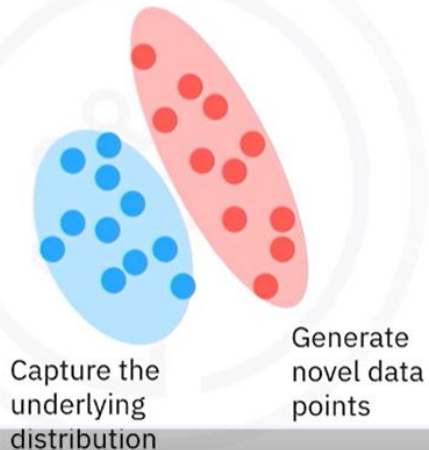
- Cannot understand context
- Cannot generate new content

Each data point labeled with its class



## Generative AI models

Create new content based on training data



## Generative AI

### Prompt

- Text
- Image
- Video
- Other forms of input



### Output

- Text
- Images
- Audio
- Video
- Code
- Data



## Discriminative AI versus generative AI

### Discriminative AI



Is this image a drawing of a nest or an egg?

### Generative AI



Draw an image of a nest with three eggs in it.

"AI can not only boost our analytic and decision-making abilities but also heighten creativity."  
-Harvard Business Review

## How does generative AI develop creativity?



Creative skills through generative AI models

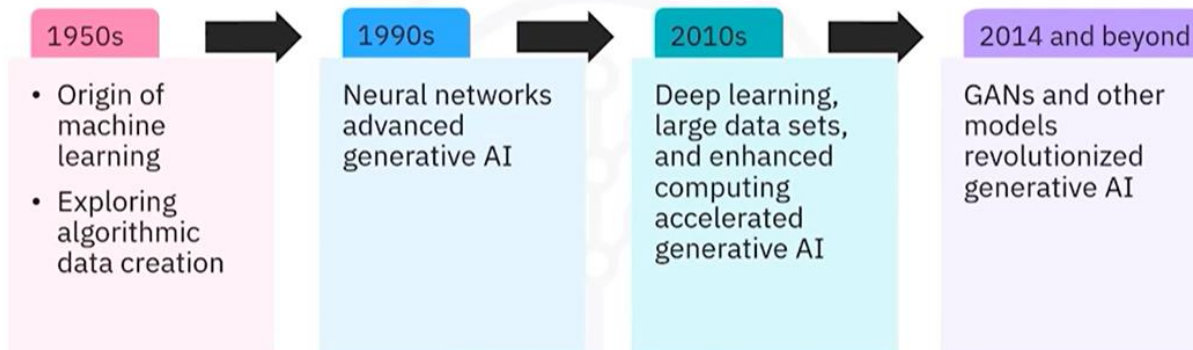
Generative adversarial networks (GANs)

Variational autoencoders (VAEs)

Transformers

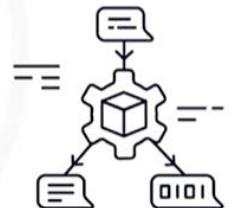
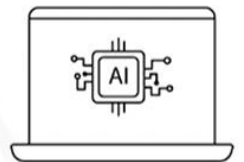
Diffusion models

## Evolution of generative AI



## Foundation models

- AI models with broad capabilities
- Adapted to build specialized and advanced models or tools
- Large language models: Process and generate text



## Foundation models: Examples

### Examples of LLMs

- **OpenAI:** GPT n-series (GPT 1, GPT2, GPT-3/3.5, and GPT-4)
- **Google:** PaLM
- **Meta:** Llama

### Examples of models for image generation

- Stable Diffusion
- DALL·E

## Recap

- Generative AI models generate new content based on the training data
- Building blocks of generative AI include GANs, VAEs, transformers, and diffusion models
- Foundation models can be adapted to create specialized models or tools
- Generative AI has potential applications in different domains and industries

## Generative AI tools

- Generative AI tools for diverse uses cases



#### Text generation

- ChatGPT
- Gemini



#### Image generation

- DALL·E 2
- Midjourney



#### Video generation

- Synthesia



#### Code generation

- Copilot
- AlphaCode



## In-context learning

- Method of prompt engineering
- Demonstrations of the task provided to the LLM as part of the prompt
- Doesn't require additional training
- New task learned from a small set of examples presented within the context at inference time



## In-context learning



### Advantages:

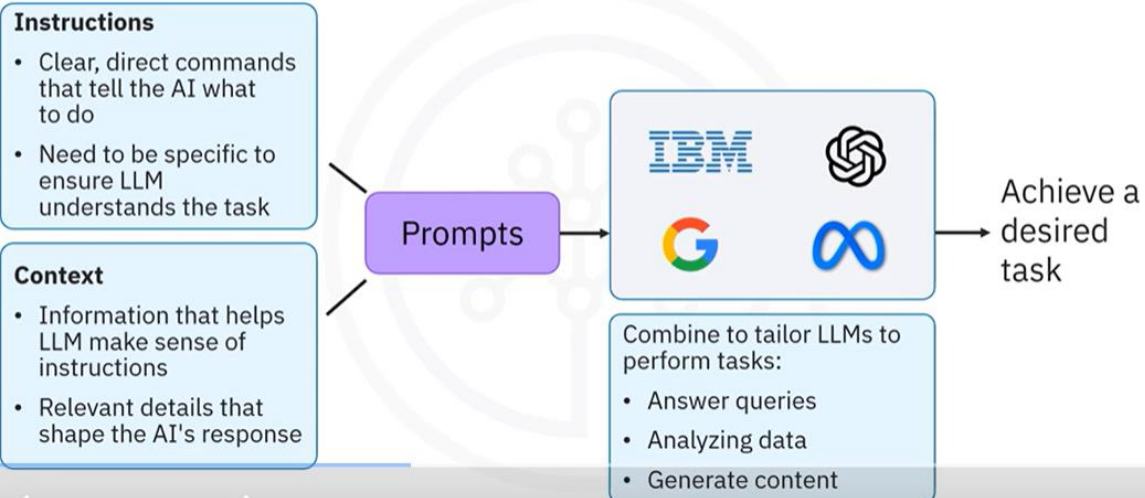
- No fine-tuning needed
- Reduces time and resource consumption
- Improves performance



### Disadvantages:

- Limited to what can fit in-context
- Complex tasks need gradient steps
- Involves adjustments based on error gradients

## What are prompts?



## What is prompt engineering?



- Designing and refining prompts to communicate with AI systems (LLMs)
- Involves crafting questions, commands, or statements to elicit accurate, relevant, and contextually appropriate responses
- Fundamental to customer service and computational linguistics

## Why prompt engineering?

Directly influences the effectiveness and accuracy of LLMs

Ensures LLMs generate relevant, precise, and contextually appropriate responses



Meets user needs through clearer prompts and reduced misunderstanding

Eliminates the need for continual fine-tuning

## First basic prompt

The wind is

Blowing gently through the trees,  
whispering secrets and stories to anyone  
who cares to listen.

## Elements of a prompt

### Instructions



Tell LLM what needs to be done

'Classify the following customer review into neutral, negative, or positive sentiment.'

### Context



Helps LLM understand the scenario

'This review is part of feedback for a recently launched product.'

### Input data



The actual data the LLM will process

'The product arrived late, but the quality exceeded my expectations.'

### Output indicator



Part of the prompt where the LLM's response is expected

'Sentiment:'

## Recap

- In-context learning: Method of prompt engineering where demonstrations of the task are provided to the model
- Prompts are inputs given to an LLM, to guide it toward performing a specific task
- Prompt engineering: Process to design and refine the questions, commands, or statements to get relevant and accurate responses
- Advantages of prompt engineering :
  - Boosts the effectiveness and accuracy of LLMs
  - Ensures relevant responses
  - Facilitates meeting user expectations
  - Eliminates the need for continual fine-tuning
- Key prompt elements: Instructions, context, input data, and output indicator



## What you will learn



Describe  
LangChain's  
purpose



Discuss its  
benefits



Devise practical  
uses for  
LangChain



Explain how  
LangChain  
works with  
other data types

## Purpose

So, what is considered sustainable or green publishing?

These terms do not yet have a clear, concise definition, but in its new stages of life, there are many various descriptors thought of and considered when looking at this type of publishing. To Dennis Stovall, director of publishing at Portland State University, "environmental and cultural" factors are the two most important concerns when discussing it. In *Rethinking Paper & Ink*, "sustainable" is defined as "using human and natural resources to meet human needs in a way that does not jeopardize the needs of future generations", and is also expressed when discussing the "triple bottom line", which measures performance "based on the economic, environmental,

and social performances of a business" (5). Ultimately, Stovall also explains that this form of publishing "now, is about education, research, and best practices that both inform the production choices we make and provide the momentum for continual innovation and improvement". Essentially, what this paper aims to do then is just that - and to explore not just the definition of green or sustainable publishing, but also the efforts and practices being made under this umbrella in the midst of a major publishing shift. This will include the exploration of environmentally friendly efforts such as book materials, sustainable practices within book-making, the cultural and humanistic importance and emphasis on proper labor, green

marketing, ethics within green publishing companies, and how the internet and e-book can be considered "green" too.

In recent years we have seen a dramatic improvement in recycling efforts, and this is certainly the most widely used sustainable printing practice among both big companies and small publishing houses. Lantern Books, who have been awarded for their paper practices, have been, according to an article from the April, 2007 issue of *Publisher's Weekly*, "among the first book publishers to develop a meaningful environmentally friendly paper policy" and has "printed most of its books on 100-percent recycled paper". Chelsea Green, a leading publishing company in sustainability practices and

Complex prompts

Extensive text-based  
data

Concise summaries

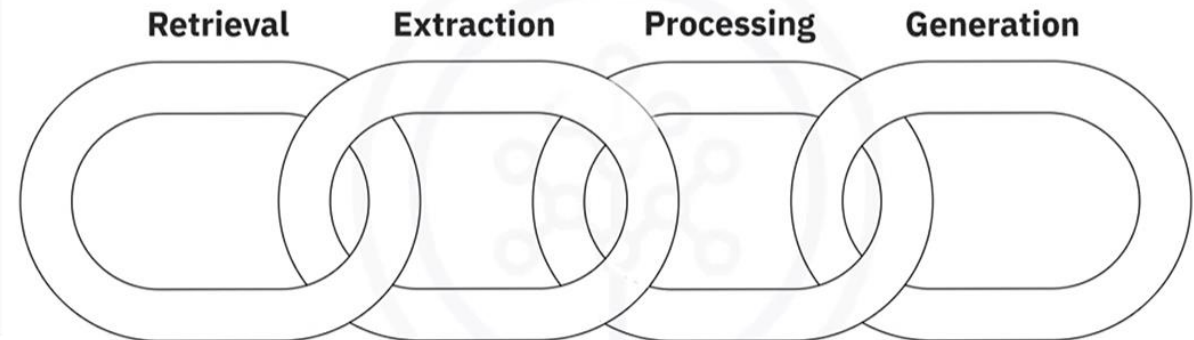
## AI Python framework



Open-source

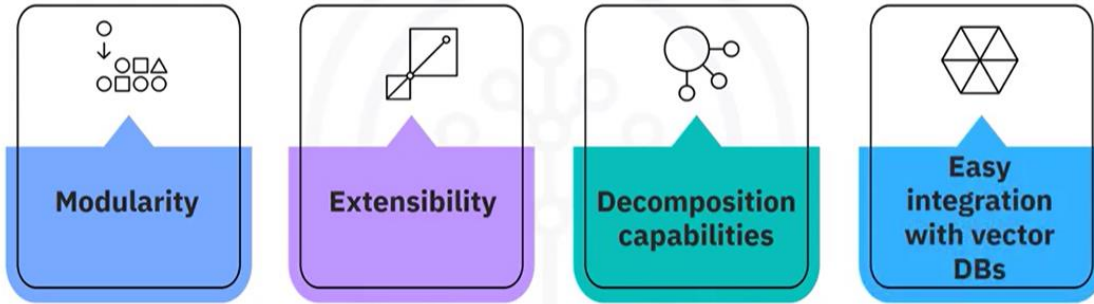
Integrate LLMs into  
AI apps

## Chaining



## Benefits

---



## Recap

---

In this video, you learned that:

- LangChain is a Python framework for pinpointing relevant information in text and providing methods for responding to complex prompts
- Benefits include modularity, extensibility, decomposition capabilities, and easy integration with vector databases
- Several practical applications include deciphering complex legal documents, extracting key statistics from reports, customer support, and automating routine writing tasks
- LangChain can be used with other data types by using external libraries and models



## Zero-shot prompt

Instructs an LLM to perform a task without any prior specific training or examples.

### Example:

```
Classify the following statement as true or false:  
'The Eiffel Tower is located in Berlin.'
```

Answer:

## Few-shot prompt

Learns from a small set of examples before tackling a similar task to generalize from a few instances to new data.

### Example:

```
Here are few examples of classifying emotions in statements:  
Statement: 'I just won my first marathon!'  
Emotion: Joy  
Statement: 'I can't believe I lost my keys again.'  
Emotion: Frustration  
Statement: 'My best friend is moving to another country.'  
Emotion: Sadness  
Now, classify the emotion in the following statement:  
Statement: 'That movie was so scary I had to cover my eyes.'  
Emotion:
```

## One-shot prompt

Gives the LLM a single example to help it perform a similar task.

### Example:

```
Translating a sentence from English to French:  
English: "How is the weather today?"  
French: "Comment est le temps aujourd'hui?"  
Now, translate the following sentence from English to French:  
English: "Where is the nearest supermarket?"  
French:
```

## Chain-of-thought (CoT) prompting



# Chain-of-thought (CoT) prompting

## Example:

### Model input

Consider the problem: 'A store had 22 apples. They sold 15 apples today and got a new delivery of 8 apples. How many apples are there now?'

Break down each step of your calculation

### Model output

Sure, let's break down the calculation step by step:

1. Initial number of apples: Start with the number of apples the store had initially, which is 22.
2. Apples sold: Subtract the number of apples sold today from the initial number. The store sold 15 apples, so subtract 15 from 22:  $22 - 15 = 7$
3. New delivery: Add the number of apples from the new delivery. The store received 8 new apples, so add 8 to the remaining apples:  $7 + 8 = 15$

So, after selling 15 apples and receiving a new delivery of 8 apples, the store now has 15 apples.

# Self-consistency

### Model input

When I was 6, my sister was half of my age. Now I am 70, what age is my sister?

Provide three independent calculations and explanations, then determine the most consistent result.

# Self-consistency

Used to enhance the reliability and accuracy of outputs



Generates multiple independent answers to the same question

Evaluates these to determine the most consistent result

# Self-consistency

### Model output

When you were 6, your sister was half your age, making her 3 years old. The age difference between you and your sister is  $6 - 3 = 3$  years.

Now that you are 70, your sister, being 3 years younger, is  $70 - 3 = 67$ .

You were 6 and your sister was half your age, so she was 3.

Since that time, 64 years have passed (from when you were 6 to when you are now 70).

Your sister's age would then be  $3 + 64 = 67$ .

The ratio when you were 6 (you were twice as old as your sister) means she was 3.

The consistent age gap is 3 years, as previously calculated.

Apply this age gap to your current age:  $70 - 3 = 67$ .

It's consistent that your sister is 67 years old now.



## Tools and applications

Facilitate interactions with LLMs:



Playground



LangChain



HuggingFace



AI classroom



**Interactive experimentation:**  
Allows real-time tweaking and testing of prompts



**Diverse model access:**  
Provides access to various pre-trained models suitable for different tasks and languages



**Collaboration and sharing:**  
Facilitates sharing and collaborative editing of prompts



**Performance tracking:**  
Offers tools to track changes, analyze results, and optimize prompts

## Prompt template



LangChain

- Uses prompt templates: Predefined recipes for generating effective prompts for LLMs
- Prompt templates include:
  - Instructions for the language model
  - A few-shot examples to help model understand context and expected responses
  - Specific question directed at the language model



## Prompt template

```
1 from langchain_core.prompts import PromptTemplate
2
3 prompt_template = PromptTemplate.from_template(
4     "Tell me a {adjective} joke about {content}."
5 )
6
```

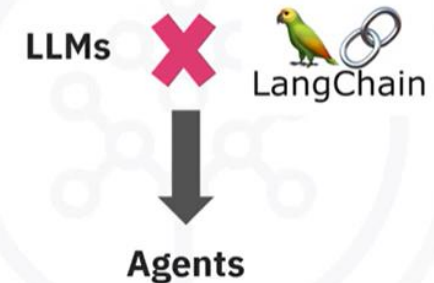
```
1 prompt_template.format(adjective="funny", content="chickens")
```



```
1 "Tell me a funny joke about chickens."
```

## Applications

LLM- and LangChain-powered agents



# Applications

---

Agents perform complex tasks across domains using different prompts

## Examples:

- Q and A agents with sources
- Content agents for creation and summarization
- Analytic agents for data analysis and business intelligence
- Multilingual agents for seamless, context-aware translation and communication



## Recap

---

- Advanced methods for prompt engineering: Zero-shot prompt, few-shot prompt, chain-of-thought prompting, and self-consistency
- Prompt engineering tools facilitate interactions with LLMs
- LangChain uses 'prompt templates,' which are predefined recipes for generating effective prompts for LLMs
- Agent: Key component in prompt applications that can perform complex tasks across various domains using different prompts





## What is LCEL?

LangChain Expression Language or LCEL

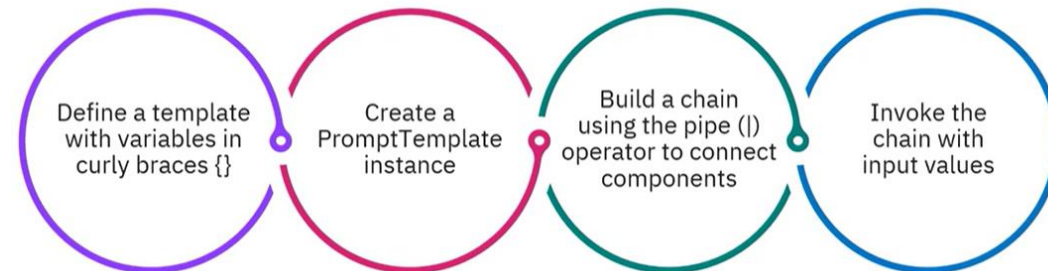
- Builds applications using the pipe (|) operator
- Ensures a clean, readable flow of data

## What is LCEL?

The new, recommended LCEL pattern provides:

- Better composability
- Clearer visualization of data flow
- Greater flexibility when constructing complex chains

## Create an LCEL pattern



## LCEL main Runnable composition primitives

### Runnables

- Interface and building blocks
- Connects the following components into a pipeline:
  - LLMs
  - Retrievers
  - Tools



## LCEL main Runnable composition primitives

### RunnableSequence

- Chains components sequentially
- Passes output from one component as input to the next

```
from langchain_core.runnables import RunnableSequence
```

```
chain = RunnableSequence([runnable1, runnable2])
```

## LCEL main Runnable composition primitives

### RunnableParallel

- Runs multiple components concurrently
- Uses the same input

```
from langchain_core.runnables import RunnableParallel
```

```
chain = RunnableParallel({  
    "key1": runnable1,  
    "key2": runnable2,  
})
```

## LCEL main Runnable composition primitives

### LCEL

- Avoids usage of RunnableSequence
- Connects runnable1 and runnable 2 with a pipe (|)

```
chain = Runnable1 | Runnable2
```

## LCEL main Runnable composition primitives

### LCEL:

Handles  
type coercion  
automatically

Converts regular  
code into  
runnable  
components

Converts  
dictionaries to  
RunnableParallel  
and functions to  
RunnableLambda

## LCEL main Runnable composition primitives

### Pipe (|) operator

- Combines the prompt templates with the LLM

### Dictionary structure

- Creates a RunnableParallel that processes all three tasks

```
translate_this_text:  
{text}) | llm,
```

```
"translation":  
ChatPromptTemplate.from_template("Transl  
ate this text to French: {text}") | llm,
```

```
"sentiment":  
ChatPromptTemplate.from_template(  

```

```
"What is the sentiment of this text?  
Answer with positive, negative, or  
neutral: {text}") | llm
```

## Creating templates and chains

- Pipe (|) operator creates a sequence by connecting runnable components
- RunnableLambda formats the prompt with variables
- Pipe (|) operator passes the formatted prompt to the LLM
- Another pipe passes the response to the StrOutputParser

```
joke_chain = (  
    RunnableLambda(format_prompt)  
    | llm  
    | StrOutputParser()  
)
```

# Run the chain

## What next?

### We covered:

- Essentials of LCEL
- Benefits and composition primitives

**LCEL:** Suited for simpler orchestration tasks

**LangGraph:** Suited for complex workflows

### LCEL's strengths:

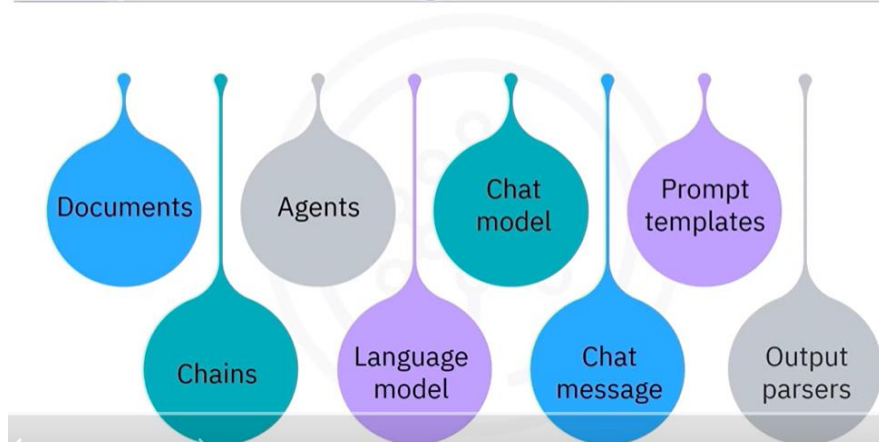
- Parallel execution
- Async support
- Simplified streaming
- Automatic tracing

## Recap

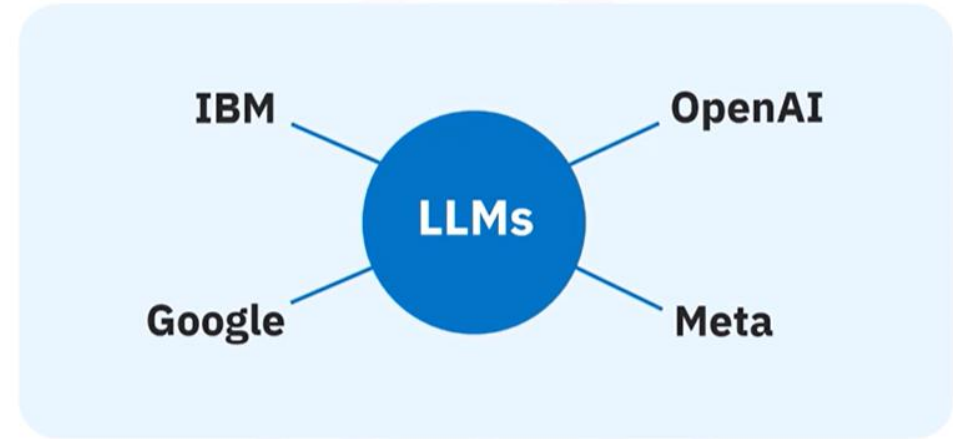
- LCEL pattern structures workflows use the pipe (|) operator
- Prompts use templates with variables in curly braces {}
- RunnableSequence links components for sequential execution
- RunnableParallel runs multiple components concurrently with the same input
- LCEL simplifies syntax by replacing RunnableSequence with the pipe operator
- Type coercion automatically converts functions and dictionaries into compatible components



## Components of LangChain



## Introduction



## Language model

```
model_id = 'mistralai/mixtral-8x7b-instruct-v01'

parameters = {
    GenParams.MAX_NEW_TOKENS: 256, # this controls the maximum number of tokens in
    the generated output
    GenParams.TEMPERATURE: 0.5, # this randomness or creativity of the model's
    responses}

credentials = {"url": "https://us-south.ml.cloud.ibm.com"}

project_id = "skills-network"

model = ModelInference (model_id=model_id, params=parameters,
credentials=credentials, project_id=project_id)
```

## Chat model

- A type of language model
- Designed for efficient conversations





# Chat model

```
model = ModelInference(.....)

mixtral_llm = WatsonxLLM(model = model)

print(mixtral_llm.invoke("Who is man's best friend?"))
```

## Chat message

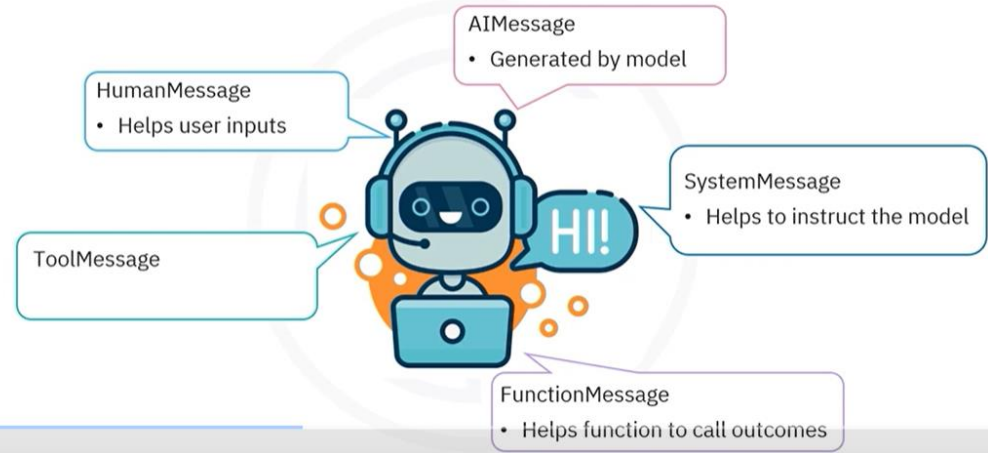
```
role    content
|       |
SystemMessage(content="You are a nice AI bot that helps a user Z
figure out what to eat in one short sentence")
```

## Chat message

```
msg = mixtral_llm.invoke(
    [
        HumanMessage(content="What month follows June?")
    ]
)

print(msg)
```

## Chat message



## Chat message

```
msg = mixtral_llm.invoke(
    [
        SystemMessage(content="You are a supportive AI bot that
suggests fitness activities to a user in one short
sentence"),
        HumanMessage(content="I like high-intensity workouts, what
should I do?"),
        AIMessage(content="You should try a CrossFit class"),
        HumanMessage(content="How often should I attend?")
    ]
)

print(msg)

AI: Aim to do CrossFit 3 times a week with rest days in between.
```

## Prompt templates



## Prompt templates

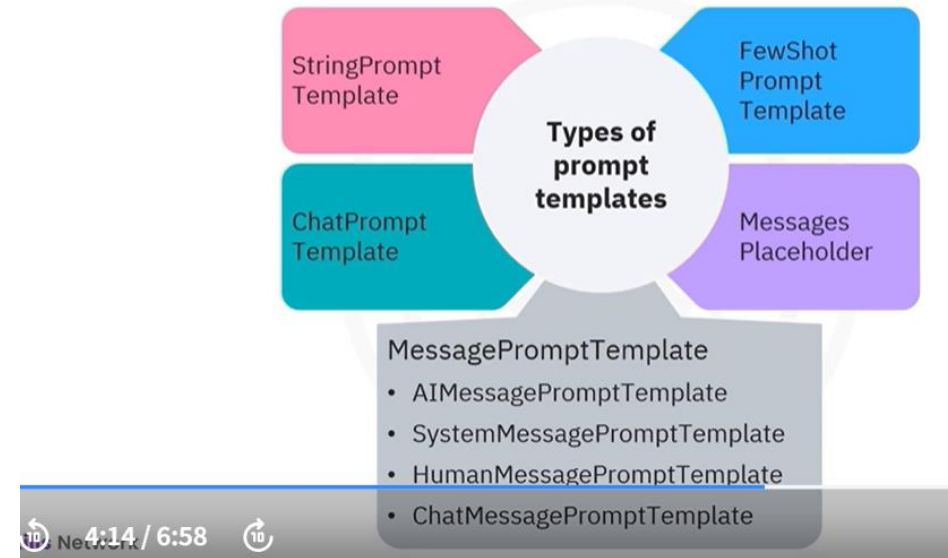
```
prompt = ChatPromptTemplate.from_messages([
    ("system", "You are a helpful assistant"),
    ("user", "Tell me a joke about {topic}")
])
```

```
input_ = {"topic": "cats"}
```

```
prompt.invoke(input_)
```

```
ChatPromptValue(messages=[SystemMessage(content='You are a helpful assistant'), HumanMessage(content='Tell me a joke about cats')])
```

## Prompt templates



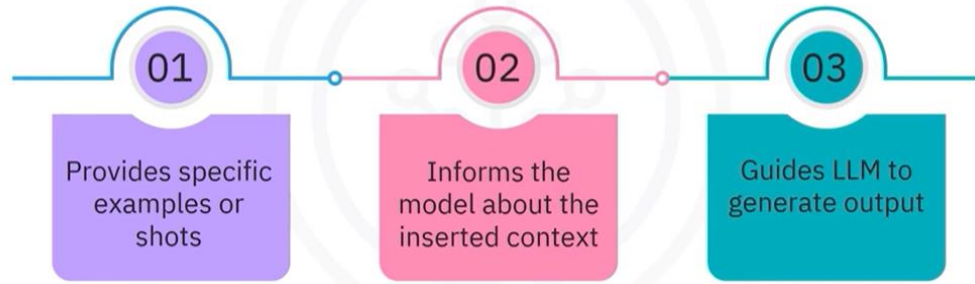
## Prompt templates: Example selector

Important:  
Selecting relevant  
examples and  
putting them into the  
prompt

Handles example  
section process  
efficiently

## Prompt templates: Example selector

- For example: FewShotPromptTemplate



## Prompt templates: Example selector

```
From langchain.prompts import PromptTemplate, FewShotPromptTemplate
```

```
From langchain.prompts.example_selector import NGramOverlapExampleSelector
```

```
NGramOverlapExampleSelector (
    examples=examples,
    example_prompt=example_prompt,
    threshold=-0.1
)
dynamic_prompt = FewShotPromptTemplate(
    example_selector=example_selector,
    example_prompt=example_prompt,
    prefix="Give the location an item is
usually found in",
    suffix="Input: {item}\nOutput:",
    input_variables=["item"],
)
```

Forward 10s (right)

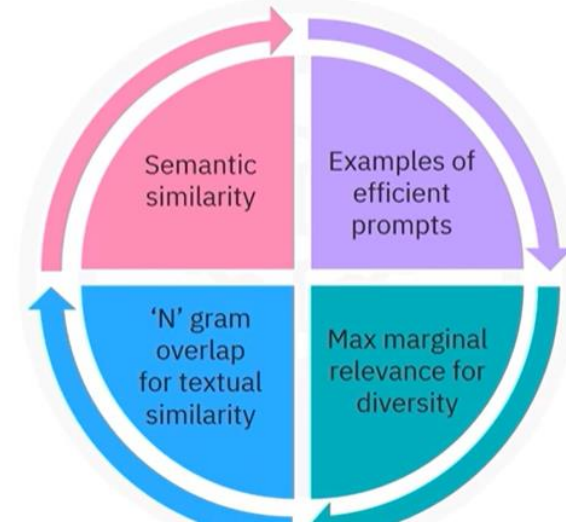
Give the location an item is  
usually found in

Example Input: pirate  
Example Output: ship

Example Input: pilot  
Example Output: plane

Input: plant  
Output:

## Prompt templates: Example selector



## Output parsers





## Output parsers

```
from langchain.output_parsers import CommaSeparatedListOutputParser

output_parser = CommaSeparatedListOutputParser()

format_instructions = output_parser.get_format_instructions()
prompt = PromptTemplate(
    template="Answer the user query. {format_instructions}\nlist five {subject}.",
    input_variables=["subject"],
    partial_variables={"format_instructions": format_instructions},
)

chain = prompt | mixtral_llm | output_parser
chain.invoke({"subject": "ice cream flavors"})

['Chocolate', 'Vanilla', 'Strawberry', 'Mint Chocolate Chip', 'Butter Pecan']
```

## Recap

- Prompt templates: Translates user's query to clear instructions
- Example selector: Informs model about the input context
- Output parsers: Transform the output from LLM to a suitable format

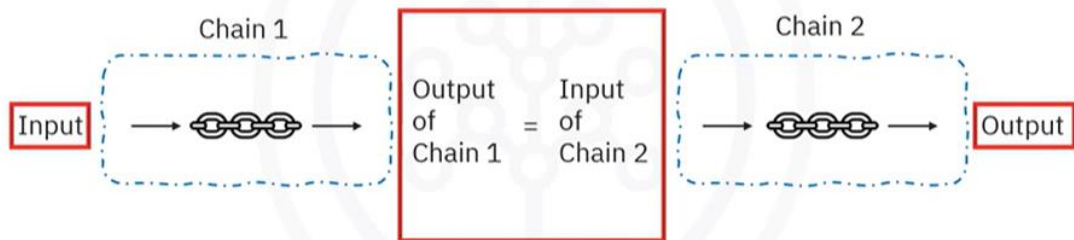
## Recap

- LangChain is an open-source interface
- Core components of LangChain
  - Language model: Foundation of LLMs
  - Chat models: Designed for efficient conversations
  - Chat messages: Handled by chat models
    - HumanMessage
    - AIMessage
    - SystemMessage
    - FunctionMessage
    - ToolMessage



# Chains

Refer to sequences of calls, whether to an LLM, a tool, or a data preprocessing step



# Chains

```
from langchain.chains import LLMChain, SequentialChain
```

```
template = """Your job is to come up with a classic dish from the  
area that the users suggests.  
{location}  
YOUR RESPONSE:
```

```
"""
```

```
prompt_template = PromptTemplate(template=template,  
input_variables=['location'])
```

```
# chain 1
```

```
location_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template,  
output_key='meal')
```

# Chains

```
template = """Given a meal {meal}, give a short and simple recipe on how  
to make that dish at home.
```

```
YOUR RESPONSE:
```

```
"""
```

```
prompt_template = PromptTemplate(template=template,  
input_variables=['meal'])
```

```
# chain 2
```

```
dish_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template,  
output_key='recipe')
```

# Chains

```
template = """Given the recipe {recipe}, estimate how much time I need to  
cook it.
```

```
YOUR RESPONSE:
```

```
"""
```

```
prompt_template = PromptTemplate(template=template,  
input_variables=['recipe'])
```

```
# chain 3
```

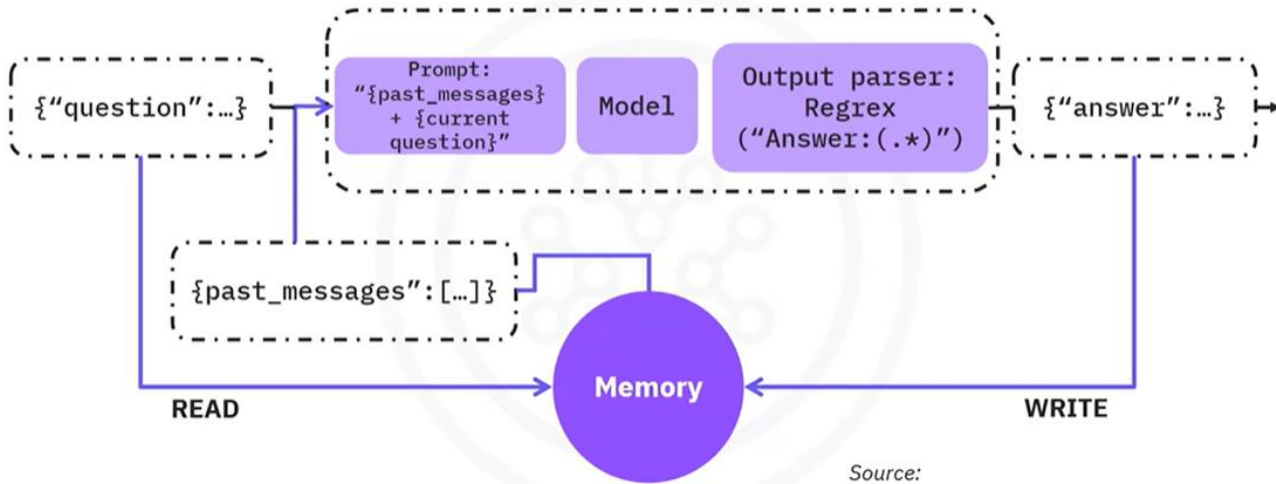
```
recipe_chain = LLMChain(llm=mixtral_llm, prompt=prompt_template,  
output_key='time')
```

## Chains

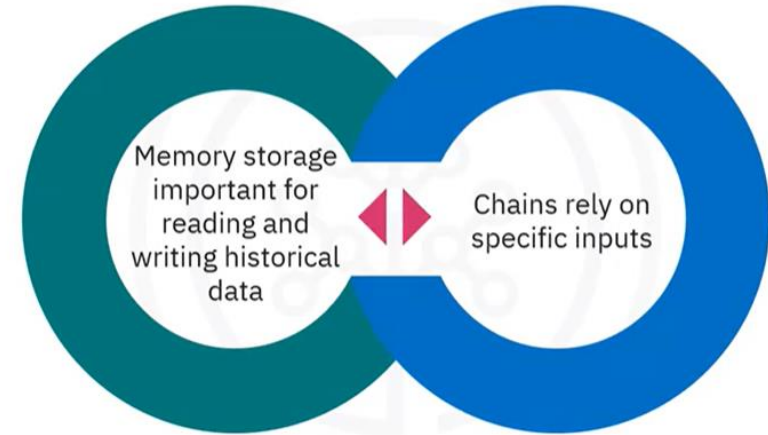
```
# overall chain
overall_chain = SequentialChain(chains=[location_chain, dish_chain,
                                     recipe_chain],
                               input_variables=['location'],
                               output_variables=['meal', 'recipe', 'time'],
                               verbose= True)

overall_chain.invoke(input={'location': 'China'})
```

## Memory



## Memory



## Memory

```
from langchain.memory import ChatMessageHistory

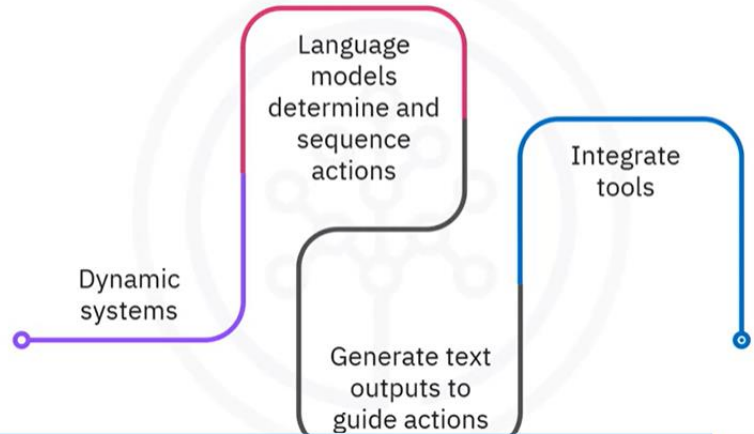
chat = mixtral_llm

history = ChatMessageHistory()
history.add_ai_message("hi!")
history.add_user_message("what is the capital of France?")
```

Memory

```
[AIMessage(content='hi!'),
 HumanMessage(content='what is the capital of France?')]
```

## Agents



## Agents



## Agents

```
df = pd.read_csv(
    "example.csv"
)

agent = create_pandas_dataframe_agent(
    mixtral_llm,
    df,
    verbose=True,
    return_intermediate_steps=True
)

agent.invoke("How many rows in the dataframe?")
```

There are 139 rows in the dataframe.

## Recap

- LangChain is a platform for embedding APIs
- Chains
  - Sequences of calls in LangChain
  - Output from step 1 becomes the input for step 2
  - Defines the template string → Creates prompt templates → Creates LLMChain object name
- Memory
  - Important for reading and writing historical data
- Agents
  - Dynamic systems where the language model determines the sequence of actions
  - Integrate with tools